


89% of AI agents
never leave the
pilot.

38% pilot. 11% ship.



● PILOTING • 38% ● IN PRODUCTION • 11%

DELOITTE TECH TRENDS 2026

Not one of them is the model.



64%

Evaluation gaps



57%

Governance friction




51%

Model reliability

Every one is a harness problem — and every harness problem traces back to an input nobody wrote down.

The brief



BECOMES

System prompt + permission boundary

The journey map



BECOMES

Tool allowlist + escalation gate

The measurement plan



BECOMES

Eval harness +
golden-set CI

The segmentation



BECOMES

Retrieval strategy + recall target


Same model. Different harness.

27.8%



EMBEDDING BASELINE

49.6%



AGENTIC HARNESS

MICROSOFT • ARXIV 2605.05538

Receipts.

DELOITTE TECH TRENDS 2026

38% piloting vs 11% in production

GARTNER, JUN 2025

40%+ of agentic projects cancelled by 2027

MICROSOFT, ARXIV 2605.05538

BRIGHT recall@1 27.8% to 49.6%

PAGERDUTY, N=1,000 EXECUTIVES

171% average ROI once in production

Build an agent harness.

Full breakdown, with the animated data blocks →
ensolabs.ai/insights